

Python Programming for Data Science

Week 41, Friday

External modules: Pandas

Pandas: Series

pandas

What is pandas

pandas

What is pandas

- A library for data manipulation and analysis.

pandas

What is pandas

- A library for data manipulation and analysis.
- Make it easy to work with structured tables of numbers

pandas

What is pandas

- A library for data manipulation and analysis.
- Make it easy to work with structured tables of numbers
- Much faster than writing loops in Python

pandas

What is pandas

- A library for data manipulation and analysis.
- Make it easy to work with structured tables of numbers
- Much faster than writing loops in Python

In short: whenever you have a list of numbers, consider using numpy **and pandas**.

pandas

What is pandas

- A library for data manipulation and analysis.
- Make it easy to work with structured tables of numbers
- Much faster than writing loops in Python

In short: whenever you have a list of numbers, consider using numpy **and pandas**.

Main difference from numpy: data is labeled.

pandas - importing

```
import pandas
```

pandas - importing

```
import pandas
```

or:

```
import pandas as pd
```

pandas: two data types

Pandas is centered around two data types

Series

Like a 1D numpy array - but with labels

Dataframe

A 2D data structure with labeled columns

pandas - Series

A Series is basically a numpy array with an index which defines labels for each entry

```
np_array = np.array([1.0, 2.0, 3.0])

# convert numpy array to pandas series
pd_series = pd.Series(np_array, index=['a', 'b', 'c'])

print(pd_series)
```

```
a    1.0
b    2.0
c    3.0
dtype: float64
```

output

pandas - Series - from a dict

You can also create a Series from a dict

```
# Convert dictionary to pandas series  
pd_series = pd.Series({'a':1.0, 'b':2.0, 'c':3.0})  
print(pd_series)
```

```
a    1.0  
b    2.0  
c    3.0  
dtype: float64
```

output

pandas - Series - behave as arrays and dicts

Series behave similar to numpy arrays and dicts

```
# Just like numpy arrays, I can operate on all elements at once  
# ... but note how the labels stay aligned  
print(pd_series * 2)
```

```
a    2.0  
b    4.0  
c    6.0  
dtype: float64
```

output

pandas - Series - behave as arrays and dicts

Series behave similar to numpy arrays and dicts

```
# Just like numpy arrays, I can operate on all elements at once  
# ... but note how the labels stay aligned  
print(pd_series * 2)
```

```
a    2.0  
b    4.0  
c    6.0  
dtype: float64
```

output

```
print(np.log(pd_series))    # I can also use functions from numpy
```

```
a    0.000000  
b    0.693147  
c    1.098612  
dtype: float64
```

output

pandas - Series - behave as arrays and dicts

Series behave similar to numpy arrays and dicts

```
# Just like numpy arrays, I can operate on all elements at once  
# ... but note how the labels stay aligned  
print(pd_series * 2)
```

```
a    2.0  
b    4.0  
c    6.0  
dtype: float64
```

output

```
print(np.log(pd_series))    # I can also use functions from numpy
```

```
a    0.000000  
b    0.693147  
c    1.098612  
dtype: float64
```

output

```
# Just like dicts, I can ask for a specific key  
print(pd_series['b'])
```

```
2.0
```

output

pandas - Series - difference to numpy

Labels are used for alignment

```
print(pd_series + pd_series[1:])
```

```
a      NaN  
b      4.0  
c      6.0  
dtype: float64
```

output

Very powerful! - no worrying about adding pears to apples

pandas - Series - difference to numpy

Labels are used for alignment

```
print(pd_series + pd_series[1:])
```

```
a      NaN  
b      4.0  
c      6.0  
dtype: float64
```

output

Very powerful! - no worrying about adding pears to apples

Note that all indices are retained - unless you explicitly drop them:

```
print((pd_series + pd_series[1:]).dropna())
```

```
b      4.0  
c      6.0  
dtype: float64
```

output

pandas - accessors

Pandas provides specific support for Series of certain types, through so-called Accessors.

Data type	Accessor
String	str
Categorical	cat
Datetime	dt
Sparse	sparse

For strings, this makes the usual methods available:

```
series = pd.Series({'a':'I', 'b':'love', 'c':'python'})  
print(pd_series.str.upper())
```

```
a      I  
b    LOVE  
c  PYTHON  
dtype: object
```

output

pandas - accessors

Pandas provides specific support for Series of certain types, through so-called Accessors.

Data type	Accessor
String	str
Categorical	cat
Datetime	dt
Sparse	sparse

For strings, this makes the usual methods available:

```
series = pd.Series({'a':'I', 'b':'love', 'c':'python'})  
print(pd_series.str.upper())
```

```
a      I  
b    LOVE  
c  PYTHON  
dtype: object
```

output

You can explicitly convert a series into strings using `.astype('string')`

pandas - Categorical

In addition to the types known from numpy, pandas also has a categorical type

Can be created in different ways:

```
series = pd.Series(['male', 'female', 'female'], dtype="category")
series = pd.Series(['male', 'female', 'female']).astype('category')
series = pd.Series(pd.Categorical(['male', 'female', 'female']))
series = pd.Series(pd.Categorical(['male', 'female', 'female'],
                                  categories=['female', 'male']))
```

Category specific functionality are accessible through `.cat`:

```
print(series.cat.categories)
```

```
Index(['female', 'male'], dtype='object')
```

output

pandas - Series - Exercise

1. Use range to create a list with values from 0 to 99 and use it to initialize a Series
2. Convert the series to type string, and calculate the length of the entries

pandas - Series - Exercise - solution

1. Use range to create a list with values from 0 to 99 and use it to initialize a Series

```
# Create range series  
series = pd.Series(range(100))
```

2. Convert the series to type string, and calculate the length of the entries

```
# Convert to string and calculate lengths  
series.astype('string').str.len()
```

Pandas: Dataframes

pandas - Dataframe

- The main data type in pandas
- Equivalent to a SAS dataset, an R data frame or an SQL table
- Can be thought of as a dict of Series

pandas - Dataframe - from numpy array

A dataframe has two sets of labels

index: row labels

columns: column labels

Dataframes can be created from numpy arrays by providing these sets of labels

```
np_array = np.arange(6).reshape((2,3))  
df = pd.DataFrame(np_array, index=['a', 'b'], columns=['col1', 'col2', 'col3'])  
print(df)
```

	col1	col2	col3
a	0	1	2
b	3	4	5

output

pandas - Dataframe - from dictionary of lists

You can also initialize from a dictionary of lists

```
dict_of_lists = {'col1': [0,3], 'col2': [1,4], 'col3': [2,5]}  
df = pd.DataFrame(dict_of_lists, index=['a', 'b'])  
print(df)
```

	col1	col2	col3
a	0	1	2
b	3	4	5

output

pandas - Dataframe - from dictionary of lists

You can also initialize from a dictionary of lists

```
dict_of_lists = {'col1': [0,3], 'col2': [1,4], 'col3': [2,5]}  
df = pd.DataFrame(dict_of_lists, index=['a', 'b'])  
print(df)
```

	col1	col2	col3
a	0	1	2
b	3	4	5

output

If index is not specified, it will use [0,1, ...]

```
df = pd.DataFrame(dict_of_lists)  
print(df)
```

	col1	col2	col3
0	0	1	2
1	3	4	5

output

pandas - Dataframe - from list of dictionaries

You can also initialize from a lists of dictionaries

```
list_of_dicts = [{'col1': 0, 'col2': 1, 'col3': 2},  
                  {'col1': 3, 'col2': 4, 'col3': 5}]  
df = pd.DataFrame(list_of_dicts, index=['a', 'b'])  
print(df)
```

	col1	col2	col3
a	0	1	2
b	3	4	5

output

pandas - Dataframe - from list of dictionaries

You can also initialize from a lists of dictionaries

```
list_of_dicts = [{'col1': 0, 'col2': 1, 'col3': 2},  
                 {'col1': 3, 'col2': 4, 'col3': 5}]  
df = pd.DataFrame(list_of_dicts, index=['a', 'b'])  
print(df)
```

	col1	col2	col3
a	0	1	2
b	3	4	5

output

There are many more ways to initialize DataFrames...

pandas - Dataframe: index, columns, values

The labels and values can be accessed using `.index`, `.columns`, and `.values`

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array,
                  index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])

print(df.index)
print(df.columns)
print(df.values)
```

output

```
Index(['a', 'b', 'c'], dtype='object')
Index(['col1', 'col2'], dtype='object')
[[0 1]
 [2 3]
 [4 5]]
```

pandas - Indexing into a dataframe

General - also supports slicing, etc (see later slide):

Get by label: `df.loc[row_label, col_label]`

Get by index: `df.iloc[row_index, col_index]`

pandas - Indexing into a dataframe

General - also supports slicing, etc (see later slide):

Get by label: `df.loc[row_label, col_label]`

Get by index: `df.iloc[row_index, col_index]`

Faster - for lookup of single values:

Get by label: `df.at[row_label, col_label]`

Get by index: `df.iat[row_index, col_index]`

pandas - Other ways to index into a dataframe

Indexing directly into a dataframe:

Get column by label:	<code>df[col_label]</code>	→Series
----------------------	----------------------------	---------

Slice rows by index range:	<code>df[start:end]</code>	→DataFrame
----------------------------	----------------------------	------------

Filter by boolean list/array	<code>df[bool_list]</code>	→DataFrame
------------------------------	----------------------------	------------

This can be a bit confusing - better to stick with `loc`, `iloc`, `at`, `iat`

pandas - .loc: selecting by labels

The .loc attribute is the primary access method

- It selects **by label!**
- Ranges are allowed: all labels between **and including** both endpoints are included

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])

print(df)

# note that both endpoints are included
print(df.loc['a':'b'])
```

output

	col1	col2
a	0	1
b	2	3
c	4	5

	col1	col2
a	0	1
b	2	3

pandas - .loc: selecting by labels (2)

You can select along both rows and columns with .loc

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])

print(df)

print(df.loc['a':'b', : 'col1'])
```

output

	col1	col2
a	0	1
b	2	3
c	4	5

	col1
a	0
b	2

pandas - .loc: selecting by labels (2)

You can select along both rows and columns with .loc

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])

print(df)

print(df.loc['a':'b', : 'col1'])
```

output

	col1	col2
a	0	1
b	2	3
c	4	5

	col1
a	0
b	2

Q: What would happen if I removed the ':' before 'col1'?

pandas - .loc: selecting by labels (2)

You can select along both rows and columns with .loc

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])

print(df)

print(df.loc['a':'b', : 'col1'])
```

output

	col1	col2
a	0	1
b	2	3
c	4	5

	col1
a	0
b	2

Q: What would happen if I removed the ':' before 'col1'?

A: It returns a series

pandas - .loc: selecting with a boolean array

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])
print(df)

print(df.loc[:, 'col1']>2)

df.loc[df.loc[:, 'col1']>2] = 0
print(df)
```

output

	col1	col2
a	0	1
b	2	3
c	4	5

a	False
b	False
c	True

Name: col1, dtype: bool

	col1	col2
a	0	1
b	2	3
c	0	0

pandas - reading/writing in various formats

`read_table()`

`read_csv()` `to_csv()`

`read_html()` `to_html()`

`read_json()` `to_json()`

`read_hdf()` `to_hdf()`

`read_excel()` `to_excel()`

`read_sas()`

`read_sql()` `to_sql()`

...

pandas - Exercise 2

1. Download:
<https://wouterboomsma.github.io/ppds2021/data/british-english>
2. Read it into a DataFrame (please use the `keep_default_na=False` option to `read_table`)
3. Figure out how to assign a different column name - e.g. 'words'
4. Figure out how to select all rows starting with the letter A

pandas - Exercise 2 - solution

2. Read it into a DataFrame

```
df = pd.read_table('british-english', keep_default_na=False, header=None)
```

3. Figure out how to assign a different column name

```
df.columns = ['words']
```

4. Figure out how to select all rows starting with the letter A

```
df.loc[df['words'].str.startswith('A')]
```

pandas - DataFrames - Adding columns

Add a new column by assigning to the relevant label

```
np_array = np.arange(6).reshape((3,2))  
df = pd.DataFrame(np_array, index=['a', 'b', 'c'], columns=['col1', 'col2'])  
df.loc[:, 'col3'] = df.loc[:, 'col1']    # Adding 'col3' - copy of `col1`
```

...or use the `.assign()` method

```
df.assign(col4 = df.loc[:, 'col1'])
```

...or `.insert()` (inserts at a specific location)

```
df.insert(0, 'col0', df.loc[:, 'col1'])
```

pandas - DataFrames - Adding columns (2)

If columns don't have the expected index, missing elements will be set to NaN

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array, index=['a', 'b', 'c'], columns=['col1', 'col2'])
print(df)
df.loc[:, 'col3'] = df.loc['b':, 'col1']    # skipping first row
print(df)
```

	col1	col2
a	0	1
b	2	3
c	4	5

	col1	col2	col3
a	0	1	NaN
b	2	3	2.0
c	4	5	4.0

output

pandas - DataFrames - Adding rows

Add a new row by assigning to the relevant label

```
np_array = np.arange(6).reshape((3,2))  
df = pd.DataFrame(np_array, index=['a', 'b', 'c'], columns=['col1', 'col2'])  
df.loc['d',:] = df.loc['a',:]    # Adding 'd' - copy of 'a'
```

...or by using the `.append()` method

```
np_array = np.arange(6).reshape((3,2))  
df = pd.DataFrame(np_array, index=['a', 'b', 'c'], columns=['col1', 'col2'])  
df.append(df.loc['a',:])    # Adding copy of 'a' - under the same name!
```

pandas - DataFrames - Adding rows

Add a new row by assigning to the relevant label

```
np_array = np.arange(6).reshape((3,2))  
df = pd.DataFrame(np_array, index=['a', 'b', 'c'], columns=['col1', 'col2'])  
df.loc['d',:] = df.loc['a',:]    # Adding 'd' - copy of 'a'
```

...or by using the `.append()` method

```
np_array = np.arange(6).reshape((3,2))  
df = pd.DataFrame(np_array, index=['a', 'b', 'c'], columns=['col1', 'col2'])  
df.append(df.loc['a',:])    # Adding copy of 'a' - under the same name!
```

Duplicate index values are allowed, but makes lookups slower.

pandas - DataFrames - data alignment

When operating with two dataframes, pandas aligns on both row and column labels

```
np_array = np.arange(9).reshape((3,3))
df1 = pd.DataFrame(np_array)
df2 = pd.DataFrame(np_array[:2,:2])
print(df1)

print(df2)

print(df1+df2)
```

output

	0	1	2
0	0	1	2
1	3	4	5
2	6	7	8

	0	1
0	0	1
1	3	4

	0	1	2
0	0.0	2.0	NaN
1	6.0	8.0	NaN
2	NaN	NaN	NaN

pandas - DataFrames - descriptive statistics

Long list of standard operations: `.mean()`, `.std()`,
`.var()`, `.min()`, `.max()`, `.sumsum()`, `.sumprod()`, ...

These functions generally take an `axis` argument, and a `skipna` argument (which default to `True`)

Missing values

As we've seen, pandas will insert NaN for missing values

Q: How are these dealt with during calculations?

Missing values

As we've seen, pandas will insert NaN for missing values

Q: How are these dealt with during calculations?

A: Unlike in numpy, pandas takes missing values into account

```
np_array = np.arange(9).reshape((3,3))
df1 = pd.DataFrame(np_array)
df2 = pd.DataFrame(np_array[:2,:2])
print(df1+df2)

print((df1+df2).sum(axis=0))
```

output

	0	1	2
0	0.0	2.0	NaN
1	6.0	8.0	NaN
2	NaN	NaN	NaN

0	6.0
1	10.0
2	0.0

dtype: float64

Missing values

As we've seen, pandas will insert NaN for missing values

Q: How are these dealt with during calculations?

A: Unlike in numpy, pandas takes missing values into account

```
np_array = np.arange(9).reshape((3,3))
df1 = pd.DataFrame(np_array)
df2 = pd.DataFrame(np_array[:2,:2])
print(df1+df2)

print((df1+df2).sum(axis=0))
```

output

	0	1	2
0	0.0	2.0	NaN
1	6.0	8.0	NaN
2	NaN	NaN	NaN

0	6.0
1	10.0
2	0.0

dtype: float64

You can also specify the minimum number of non-NaN values, using the `min_count` option.

Missing values (2)

You can detect missing values using `.isna()` and `.notna()`

```
np_array = np.arange(9).reshape((3,3))
df1 = pd.DataFrame(np_array)
df2 = pd.DataFrame(np_array[:2,:2])
print(df1+df2)

print((df1+df2).isna())
```

output

	0	1	2
0	0.0	2.0	NaN
1	6.0	8.0	NaN
2	NaN	NaN	NaN

	0	1	2
0	False	False	True
1	False	False	True
2	True	True	True

Missing values (3)

You can replace missing values using `.fillna()`

```
np_array = np.arange(9).reshape((3,3))
df1 = pd.DataFrame(np_array)
df2 = pd.DataFrame(np_array[:2,:2])
print(df1+df2)

print((df1+df2).fillna(0.0))
```

output

	0	1	2
0	0.0	2.0	NaN
1	6.0	8.0	NaN
2	NaN	NaN	NaN

	0	1	2
0	0.0	2.0	0.0
1	6.0	8.0	0.0
2	0.0	0.0	0.0

Missing values (4)

You can also:

- Drop missing values using `.dropna()`
- Replace missing by interpolation `.interpolate()`
- ...

pandas - sorting labels

`.sort_index()` sorts either by row or column labels:

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array,
                   index=['a', 'b', 'c'],
                   columns=['col1', 'col2'])
print(df.sort_index(axis=1, ascending=False))
```

output

	col2	col1
a	1	0
b	3	2
c	5	4

pandas - sorting labels

`.sort_index()` sorts either by row or column labels:

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array,
                   index=['a', 'b', 'c'],
                   columns=['col1', 'col2'])
print(df.sort_index(axis=1, ascending=False))
```

output

	col2	col1
a	1	0
b	3	2
c	5	4

There is an `inplace` option which changes the original instead of returning a new data frame

pandas - sorting values

`.sort_values(by=label)` sorts either by row or column labels:

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array,
                  index=['a', 'b', 'c'],
                  columns=['col1', 'col2'])
print(df.sort_values(by='col2', ascending=False))
```

output

	col1	col2
c	4	5
b	2	3
a	0	1

pandas - sorting values

`.sort_values(by=label)` sorts either by row or column labels:

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array,
                   index=['a', 'b', 'c'],
                   columns=['col1', 'col2'])
print(df.sort_values(by='col2', ascending=False))
```

output

	col1	col2
c	4	5
b	2	3
a	0	1

There is an `inplace` option which changes the original instead of returning a new data frame

pandas - sorting values

`.sort_values(by=label)` sorts either by row or column labels:

```
np_array = np.arange(6).reshape((3,2))
df = pd.DataFrame(np_array,
                   index=['a', 'b', 'c'],
                   columns=['col1', 'col2'])
print(df.sort_values(by='col2', ascending=False))
```

output

	col1	col2
c	4	5
b	2	3
a	0	1

There is an `inplace` option which changes the original instead of returning a new data frame

You can sort by multiple labels by providing a list to `by=`

pandas - DataFrames - concatenating

```
np_array = np.arange(6).reshape((3,2))
df1 = pd.DataFrame(np_array, index=['a', 'b', 'c'],
                    columns=['col1', 'col2'])

print(df1)

df2 = pd.DataFrame(np_array[:2, :1],
                    index=['a', 'b'],
                    columns=['col3', ])

print(df2)

print(pd.concat([df1, df2], axis=1))
```

output

	col1	col2
a	0	1
b	2	3
c	4	5

	col3
a	0
b	2

	col1	col2	col3
a	0	1	0.0
b	2	3	2.0
c	4	5	NaN

pandas - DataFrames - merging

`.merge()` is an efficient implementation of SQL-like merge operations

```
df1 = pd.DataFrame({'id' : [1, 2, 3],  
                    'gender' : pd.Categorical(['male', 'female', 'female'])})  
df2 = pd.DataFrame({'id' : [1, 2, 3],  
                    'name' : pd.Categorical(['Bob', 'Alice', 'Anna'])})  
pd.merge(df1, df2, on='id')
```

	gender	id	name
0	male	1	Bob
1	female	2	Alice
2	female	3	Anna

output

You can also merge on the index labels

```
df1 = pd.DataFrame({'gender' : pd.Categorical(['male', 'female', 'female'])})  
df2 = pd.DataFrame({'name' : pd.Categorical(['Bob', 'Alice', 'Anna'])})  
pd.merge(df1, df2, left_index=True, right_index=True)
```

pandas - DataFrames - grouping

The `.groupby()` groups by value

```
df = pd.DataFrame({'name' : pd.Categorical(['Bob', 'Alice', 'Anna']),  
                  'gender' : pd.Categorical(['male', 'female', 'female']),  
                  'height': [170, 180, 165]})  
gender_grps = df.groupby('gender')  
print(gender_grps)
```

```
<pandas.core.groupby.DataFrameGroupBy object at 0x11861bd30>
```

output

You can also group by a dynamically created series:

```
df.groupby(df['name'].str.len())
```

pandas - DataFrames - grouping

The `.groupby()` groups by value

```
df = pd.DataFrame({'name' : pd.Categorical(['Bob', 'Alice', 'Anna']),  
                  'gender' : pd.Categorical(['male', 'female', 'female']),  
                  'height': [170, 180, 165]})  
gender_grps = df.groupby('gender')  
print(gender_grps)
```

```
<pandas.core.groupby.DataFrameGroupBy object at 0x11861bd30>
```

output

You can also group by a dynamically created series:

```
df.groupby(df['name'].str.len())
```

What can you do with these groups?

pandas - DataFrames - grouping (2)

1. The .groups attribute provides some information

```
print(gender_grps.groups)
```

output

```
{'female': Int64Index([1, 2], dtype='int64'), 'male': Int64Index([0], dtype='int64')}
```

2. You can also iterate over the groups

```
for name, grp in gender_grps:  
    print(name)  
    print(grp)
```

output

```
female  
   gender  height  name  
1  female    180  Alice  
2  female    165  Anna  
male  
   gender  height  name  
0   male    170   Bob
```


pandas - DataFrames - grouping (3)

...or you can calculate summary statistics

```
print(gender_grps.agg({'height':np.average}))
```

	height
gender	
female	172.5
male	170.0

output

pandas - grouping - Exercise

1. Download:
<https://wouterboomsma.github.io/ppds2021/data/british-english>, and read it into a dataframe
2. Get pandas to group by first letter - and use this to count the words for each letter

pandas - grouping - Exercise - solution

1. Download:

<https://wouterboomsma.github.io/ppds2021/data/british-english>, and read it into a dataframe

```
df = pd.read_table('british-english', keep_default_na=False, header=None)
```

2. Get pandas to group by first letter - and use this to count the words for each letter

```
df.groupby(df[0].str[0]).count()
```

pandas - working with time data

Pandas has been designed for use with time series data

You can create a time range using `pd.date_range()`, and use this as index in a Series

```
idx = pd.date_range('26/9/2018 08:00', periods=3, freq='H')
series = pd.Series(np.arange(len(idx)), index=idx)
print(series)
```

```
2018-09-26 08:00:00    0
2018-09-26 09:00:00    1
2018-09-26 10:00:00    2
Freq: H, dtype: int64
```

output

pandas - working with time data (2)

We can change the frequency using `.asfreq`

```
print(series.asfreq('30Min'))
```

```
2018-09-26 08:00:00    0.0  
2018-09-26 08:30:00    NaN  
2018-09-26 09:00:00    1.0  
2018-09-26 09:30:00    NaN  
2018-09-26 10:00:00    2.0
```

output

Use the `method` option to specify how to fill new values.
Alternatively:

```
print(series.asfreq('30Min').interpolate())
```

```
2018-09-26 08:00:00    0.0  
2018-09-26 08:30:00    0.5  
2018-09-26 09:00:00    1.0  
...
```

output

pandas - working with time data (2)


Too much to cover here, but you can:

- Using indexing and slicing on the time index
- Index with partial data strings
- Use `BusinessDay` as time unit - and define custom versions
- Automatically skip holidays
- Resampling time (time-based groupby)
- ...

pandas - plotting

Dataframes have associated plotting capabilities through matplotlib

```
idx = pd.date_range('26/9/2018 08:00',  
                    periods=7, freq='H')  
df = pd.DataFrame(np.arange(14).reshape([7,2]),  
                  columns=['first', 'second'],  
                  index=idx)  
df.plot()
```




Note how the data frame labels are automatically used.

pandas - plotting

Dataframes have associated plotting capabilities through matplotlib

```
idx = pd.date_range('26/9/2018 08:00',  
                    periods=7, freq='H')  
df = pd.DataFrame(np.arange(14).reshape([7,2]),  
                  columns=['first', 'second'],  
                  index=idx)  
df.plot()
```



Note how the data frame labels are automatically used.

Lots of other plot types available, using attributes under plot:
`df.plot.bar`, `df.plot.scatter`, etc

pandas - plotting - exercise

1. Create a histogram of the word counts from the previous exercise.
2. Bonus: Try to get meaningful labels on the x-axis and in the legend.

pandas - plotting - exercise - solution

1. Create a histogram of the word counts from the previous exercise.

```
# From previous exercise
df = pd.read_table('british-english', keep_default_na=False, header=None)
result = df.groupby(df.loc[:,0].str[0]).count()

# This exercise:
result = result.rename_axis(index='letter') # rename index axis
result.columns = ['word count'] # Give column a more informative name
result.plot.bar()
```